



COSCon'25

第十届中国开源年会

众智开源 | Open Source, Open Intelligence

Velox, 模块化计算引擎库在Meta的应用

Velox, Composable Compute Engine at Meta



CONTENTS

目录

- 01 Motivation**
- 02 The Velox Library**
- 03 Open Source**
- 04 Applications of Velox**





欢迎加入我们首届 **VeloxCon China 2025**

You are welcomed to join us for our first VeloxCon China 2025

12月13日（下周六）| 北京蚂蚁集团办公室

December 13 (Next Saturday) | Beijing @ Ant Group Offices



<https://veloxcon.io>

PART 01 Motivation



<https://veloxcon.io>

User Workload Variety



Transactional



Analytics



Realtime



Monitoring



Graph



ML



...?



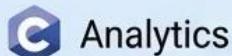
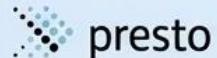
Engine Specialization



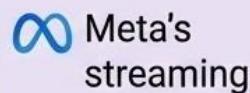
“One size does not fit all”



Analytics



Realtime Infra



Graph



Graph Database



Monitoring



Real Time Data



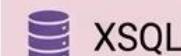
Time Series



Logging systems



Transactional



ML



Training Data



PyTorch



ML Feature Framework

Engine Specialization



- Very limited reusability.
- Duplicates efforts and forces engineers to reinvent the wheel.
- Hard to maintain and enhance.
 - Where do we optimize?
- Exposes inconsistencies to end-users.
- Hurts our capacity to move fast and innovate.



Through the Looking Glass



Different, but not really...

	Presto	Spark	DPP	XStream
Language Frontend	<i>Presto SQL</i>	<i>HQL / DataSet API</i>	<i>Koski DataFrame</i>	<i>UPM</i>
IR	<i>Presto IR</i>	<i>Spark IR</i>	<i>Koski IR</i>	<i>UPM IR</i>
Optimizer	<i>Presto Optimizer</i>	<i>Catalyst</i>	<i>Koski Optimizer</i>	<i>XStream Optimizer</i>
Execution Runtime	<i>Presto Runtime</i>	<i>Spark Runtime</i>	<i>DPP</i>	<i>XStream Runtime</i>
Execution Engine	<i>Presto</i>	<i>Spark</i>	<i>Koski</i>	<i>XStream Codegen</i>



Through the Looking Glass



Different, but not really...

	Presto	Spark	DPP	XStream
Language Frontend	Presto SQL	HQL / DataSet API	Koski DataFrame	UPM
IR	Presto IR	Spark IR	Koski IR	UPM IR
Optimizer	Presto Optimizer	Catalyst	Koski Optimizer	XStream Optimizer
Execution Runtime	Presto Runtime	Spark Runtime	DPP	XStream Runtime
Execution Engine	Presto	Spark	Koski	XStream Codegen

All very similar!



Through the Looking Glass



Different, but not really...

	Presto	Spark	DPP	XStream
Language Frontend	Presto SQL	HQL / DataSet API	Koski DataFrame	UPM
IR	Presto IR	Spark IR	Koski IR	UPM IR
Optimizer	Presto Optimizer	Catalyst	Koski Optimizer	XStream Optimizer
Execution Runtime	Presto Runtime	Spark Runtime	DPP	XStream Runtime
Execution Engine				

All very similar!

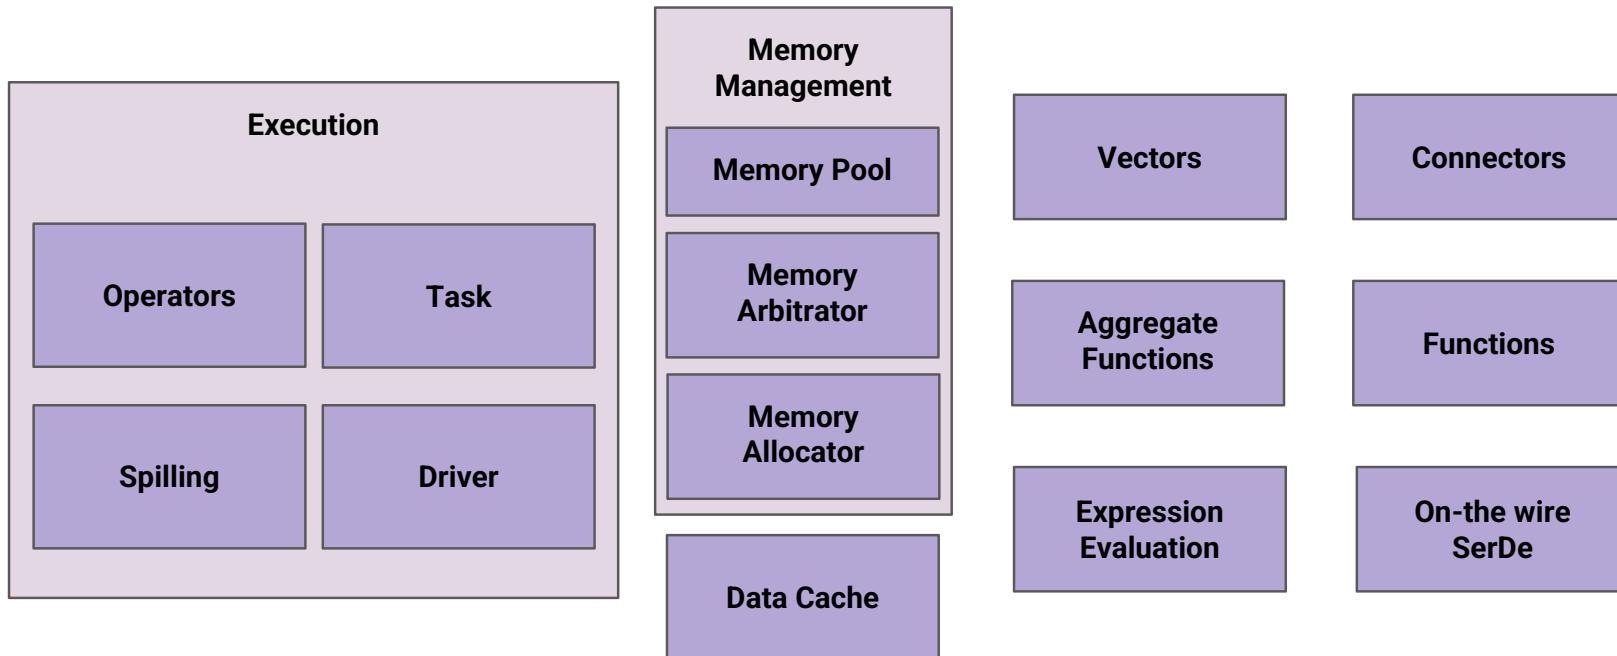
Velox



PART 02 The Velox Library



Velox Engine Library



PART 03 Open Source



Velox - Open Source



- Publicly announced [in Oct'22](#)
 - <https://engineering.fb.com/2022/08/31/open-source/velox/>
- Available in github:
 - <https://github.com/facebookincubator/velox>
- Fast growing open source community
 - +250 developers
 - Meta, IBM, Intel, Voltron Data, Microsoft, Google, ...

The screenshot shows the GitHub repository page for 'facebookincubator / velox'. The repository has 275 branches and 1 tag. The 'About' section describes Velox as a composable and fully extensible C++ execution engine library for data management systems. It includes links to the README, Apache 2.0 license, code of conduct, security policy, and archive.

VeloxCon China 2025

- Chinese communities have made tremendous contributions to Velox.
- VeloxCon China to bring the communities closer together

Engineering at Meta

Open Source ▾ Platforms ▾ Infrastructure Systems ▾ Physical Infrastructure ▾ Video Engineering & AR/VR ▾

POSTED ON AUGUST 31, 2022 TO DATA INFRASTRUCTURE, OPEN SOURCE

Introducing Velox: An open source unified execution engine



<https://veloxcon.io>



重要 Open Source 消息将会在 **VeloxCon China 2025** 公布

Important Open Source news will be announced at VeloxCon China 2025

欢迎加入我们首届 **VeloxCon China 2025**

You are welcomed to join us for our first VeloxCon China 2025

12月13日（下周六）| 北京蚂蚁集团办公室

December 13 (Next Saturday) | Beijing @ Ant Group Offices



<https://veloxcon.io>

PART 04 Applications of Velox





- First integrations:

- **Meta's XStream - Streaming System**

- Runs entirely on Velox today.

- **Meta's ETL Ingestion System**

- data warehouse and database ingestion.

- **Meta's Messaging System**

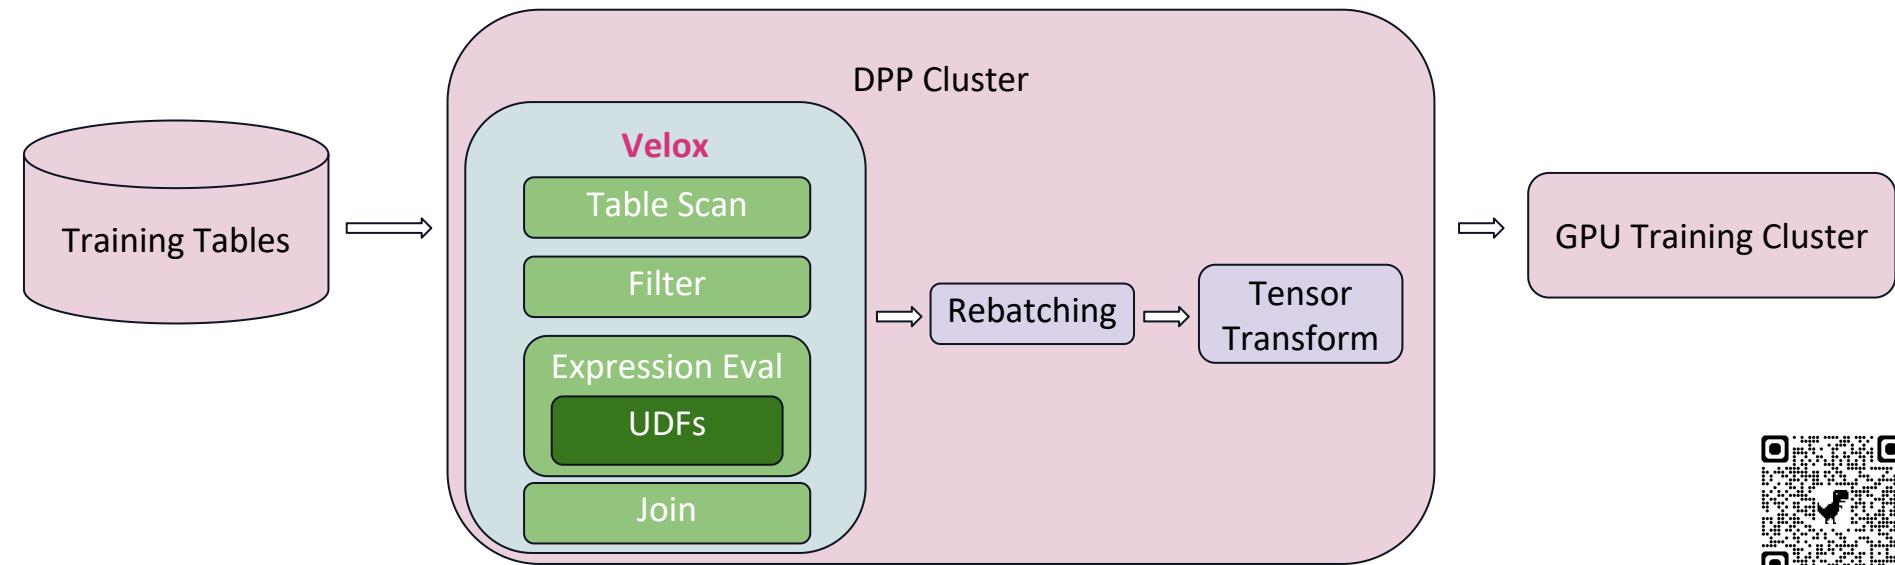


<https://veloxcon.io>

Use Cases - DPP



- DPP - Data Pre-Proc (Training data loading):
 - DPP today built on Koski
 - **Veloski** => Koski with Velox as eval.
 - Ongoing work; handful of initial models in prod using Veloski.



<https://velocon.io>

Use Cases: Presto -> Prestissimo



- Prestissimo: Presto with a Velox-based worker.
 - Same JVM coordinator.
 - Worker is C++ only (Velox).
 - Supports Vector Search
- Transparent to users - same PrestoSQL semantic.
- **3-10x** efficiency gains (depending on workload).
- Almost 100% rollout to different Presto flavors (Streaming Shuffle):
- OSS: available on IBM's watsonX.data cloud.



Use Cases - Presto-on-Spark -> Native-Presto-on-Spark



- New architecture for Presto-on-Spark
 - Same Presto-on-Spark frontend
 - Executor of modified Spark JVM with sidecar Prestissimo worker (Velox).
 - Supports Vector Search
- Transparent to users - same PrestoSQL semantic.
- **2x** efficiency gains.
- 10% already in production.
- OSS: code available on prestodb + velox



<https://veloxcon.io>

Use Cases - PySpark/Velox



● PySpark -> **PySpark/Velox**:

- Architecture overview
 - Python frontend to build Velox plan;
 - Spark scheduling to coordinate work distribution;
 - Velox backend for high efficiency compute
- Already in prod; limited query shapes.
- Focused on AI customers.



Use Cases - Apache Gluton



- **Apache Gluton:** transparent to users

- Ongoing work at Meta
- Available at Google, Microsoft, IBM, and Alibaba clouds.



<https://veloxcon.io>

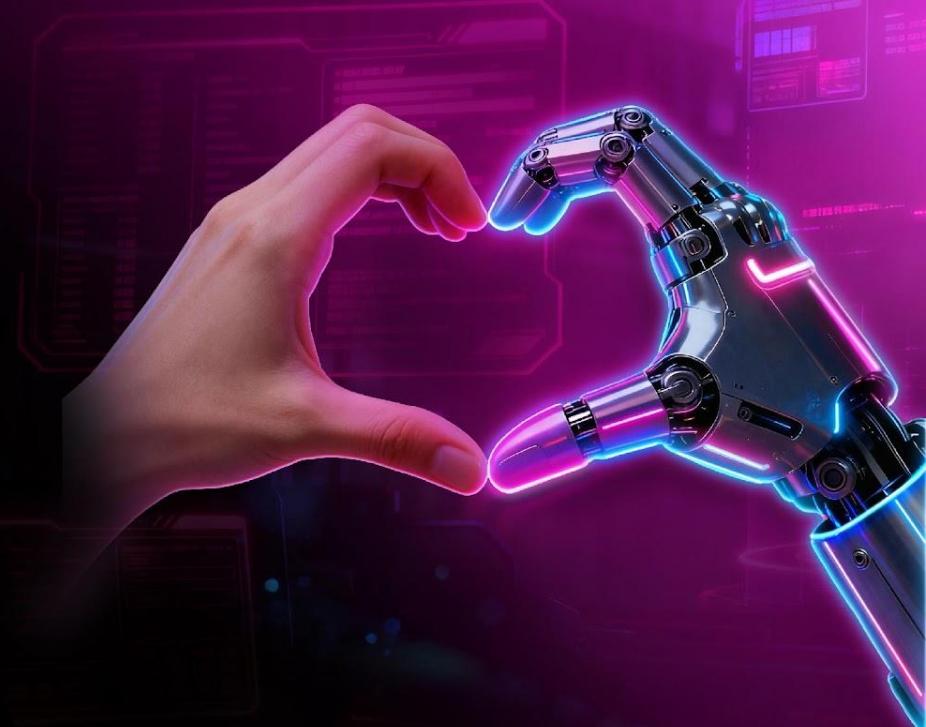


COSCon'25

第十届中国开源年会

众智开源 | Open Source, Open Intelligence

Thanks





欢迎加入我们首届 **VeloxCon China 2025**

You are welcomed to join us for our first VeloxCon China 2025

12月13日（下周六）| 北京蚂蚁集团办公室

December 13 (Next Saturday) | Beijing @ Ant Group Offices



<https://veloxcon.io>